

Tugas 1 Individu *Machine Learning*

Nama : Alfirsa Damasyifa Fauzulhaq

NIM : 226150100111019

Kelas : Machine Learning – A

Tugas

- Build a convolutional network with pre-processing on the input data (jittering, normalization). Also add dropout regularization, batch normalization, and at least one additional convolutional layer which achieves at least 90% test accuracy (for any training epoch) on MNIST dataset. Your part 1 network should train under 10 minutes, without GPUs.
- Fine-tune AlexNet to achieve at least 80% test accuracy on the MNIST dataset. Your network should train under 10 minutes, without GPUs.

Link Repository GitHub :

<https://github.com/AFauzulh/MK-Machine-Learning>

Jawab :

1) Model CNN

- *Source Code*

```
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import datasets, transforms

from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt

transform = transforms.Compose([
    transforms.Lambda(lambda image:
image.convert('RGB')),
    transforms.ToTensor(),
    # transforms.Resize((64,64)),
    transforms.ColorJitter(brightness=(0.5,1.5), contrast=(1),
saturation=(0.5,1.5), hue=(-0.1,0.1)),
    transforms.Normalize(mean=[0.485,
0.456, 0.406], std=[0.229, 0.224, 0.225])
])

train_data = datasets.MNIST(root='./Data', train=True, download=True,
transform=transform)
test_data = datasets.MNIST(root='./Data', train=False, download=True,
transform=transform)
```

```

torch.manual_seed(42)
train_loader = DataLoader(train_data, batch_size=128, shuffle=True)
test_loader = DataLoader(test_data, batch_size=128, shuffle=True)

class ConvMNIST(nn.Module):

    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 16, 3, 1)
        self.conv2 = nn.Conv2d(16, 32, 3, 1)
        self.conv2_bn = nn.BatchNorm2d(32)
        self.fc1 = nn.Linear(5*5*32, 128)
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, 10)

        self.dropout = nn.Dropout(.4)

    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2, 2)
        x = self.conv2(x)
        x = self.conv2_bn(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2, 2)

        x = x.view(-1, 32*5*5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.dropout(x)
        x = self.fc3(x)

        return F.log_softmax(x, dim=1)

device = torch.device("cpu")
model = ConvMNIST()
model.to(device)

criterion = nn.NLLLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)

import time

start_time = time.time()

epochs = 3
train_losses = []
test_losses = []
train_corrects = []
test_corrects = []

for i in range(epochs):

    train_correct = 0
    for batch, (X_train, y_train) in enumerate(train_loader):
        batch+=1

        X_train = X_train.to(device)
        y_train = y_train.to(device)

```

```

y_pred = model(X_train)
loss = criterion(y_pred, y_train)

predicted = torch.max(y_pred.data, 1)[1]
batch_correct = (predicted==y_train).sum()
train_correct+=batch_correct

optimizer.zero_grad()
loss.backward()
optimizer.step()

if batch%256 == 0:
    print(f"Epoch {i+1} Batch {batch} loss: {loss.item()}")

train_losses.append(loss)
train_corrects.append(train_correct)

test_correct = 0
with torch.no_grad():
    for batch, (X_test, y_test) in enumerate(test_loader):

        X_test = X_test.to(device)
        y_test = y_test.to(device)

        y_val = model(X_test)

        predicted = torch.max(y_val.data,1)[1]

        test_correct+=(predicted==y_test).sum()

    loss = criterion(y_val, y_test)
    test_losses.append(loss)
    test_corrects.append(test_correct)

current_time = time.time()
total = current_time - start_time
print(f"Training finished in {total/60} minutes")

for i in range(len(train_losses)):
    train_losses[i] = train_losses[i].detach().numpy()

for i in range(len(test_losses)):
    test_losses[i] = test_losses[i].detach().numpy()

plt.plot(train_losses, label='train_loss')
plt.plot(test_losses, label='test_loss')
plt.ylabel('Losses')
plt.xlabel('epoch')
plt.xticks(range(0,5))
plt.legend()

test_loader_all = DataLoader(test_data, batch_size=10000,
shuffle=False)

with torch.no_grad():
    correct = 0
    for X_test, y_test in test_loader_all:
        y_val = model(X_test)
        predicted = torch.max(y_val, 1)[1]
        correct+=(predicted==y_test).sum()

```

```
print(correct.item()/len(test_data))
print(classification_report(y_test, predicted))
print(confusion_matrix(y_test, predicted))
```

- Arsitektur Model :


```
ConvMNIST(
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
  (conv2_bn): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (fc1): Linear(in_features=800, out_features=128, bias=True)
  (fc2): Linear(in_features=128, out_features=64, bias=True)
  (fc3): Linear(in_features=64, out_features=10, bias=True)
  (dropout): Dropout(p=0.4, inplace=False)
)
```

- Proses Pelatihan

```
Epoch 1 Batch 256 loss: 0.2662428915500641
Epoch 2 Batch 256 loss: 0.05799808353185654
Epoch 3 Batch 256 loss: 0.09931761026382446
Training finished in 4.106444525718689 minutes
```

Waktu Pelatihan : ± 4.1 Menit

- Hasil Loss Pelatihan



- Hasil Akurasi Terhadap Data Test

	precision	recall	f1-score	support
0	0.98	0.98	0.98	980
1	0.98	0.99	0.98	1135
2	0.98	0.97	0.98	1032
3	0.94	0.99	0.97	1010
4	0.91	0.99	0.95	982
5	0.98	0.97	0.98	892
6	0.97	0.97	0.97	958
7	0.98	0.96	0.97	1028
8	0.99	0.92	0.96	974
9	0.98	0.95	0.97	1009
accuracy			0.97	10000
macro avg	0.97	0.97	0.97	10000
weighted avg	0.97	0.97	0.97	10000

Akurasi terhadap data uji : 97%

- Hasil *Confusion Matrix*

```
[[ 963   1   2   0   4   1   7   1   1   0]
 [   0 1120   2   4   7   0   1   1   0   0]
 [   1   5 1003   6   9   0   1   6   1   0]
 [   0   2   2 999   5   1   0   0   0   1]
 [   0   1   1   0 972   0   3   1   1   3]
 [   1   0   0  12   4 868   4   1   1   1]
 [   5   3   0   1  15   3 931   0   0   0]
 [   1   8  10   3   8   1   0 989   0   8]
 [   8   1   4  28  12   5   9   5 898   4]
 [   3   1   0   6  31   6   0   2   3 957]]
```

2) Model AlexNet + Fine Tune

Terjadi *error* size terlalu kecil pada data MNIST (28x28) saat maxpool sehingga gambar harus diresize menjadi ukuran yang lebih besar.

```
File c:\Users\HP\anaconda3\envs\torchcpu\lib\site-packages\torch\nn\functional.py:782, in _max_pool2d(input,
    780 if stride is None:
    781     stride = torch.jit.annotate(List[int], [])
--> 782 return torch.max_pool2d(input, kernel_size, stride, padding, dilation, ceil_mode)

RuntimeError: Given input size: (192x2x2). Calculated output size: (192x0x0). Output size is too small

transform = transforms.Compose([
    transforms.Resize((64,64)),
    transforms.ToTensor(),
])
```

- *Source Code*

```
import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
from torch.utils.data import DataLoader
from torchvision import datasets, transforms

from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import matplotlib.image as mpimg

transform = transforms.Compose([
    transforms.Resize((64,64)),
    transforms.ToTensor(),
])

train_data = datasets.MNIST(root='./Data', train=True, download=True,
transform=transform)
test_data = datasets.MNIST(root='./Data', train=False, download=True,
transform=transform)

image, label = train_data[0]
print(plt.imshow(image.squeeze(0)))

torch.manual_seed(42)
train_loader = DataLoader(train_data, batch_size=128, shuffle=True)
test_loader = DataLoader(test_data, batch_size=128, shuffle=True)

model = torchvision.models.alexnet(pretrained=True)
```

```

# Freeze Layers
for param in model.parameters():
    param.requires_grad = False

# Add tuning layer
model.features[0] = nn.Conv2d(1, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
model.classifier[6] = nn.Linear(4096, 10)
model.classifier.add_module('7', nn.LogSoftmax(dim=1))

for name, param in model.named_parameters():
    print(f"parameter name {name} gradient \t: {param.requires_grad}")

criterion = nn.NLLLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.0001)

device = 'cpu'
model.to(device)

import time
gc.collect()

start_time = time.time()

epochs = 3
train_losses = []
test_losses = []
train_corrects = []
test_corrects = []

for i in range(epochs):

    train_correct = 0
    for batch, (X_train, y_train) in enumerate(train_loader):
        batch+=1

        X_train = X_train.to(device)
        y_train = y_train.to(device)

        y_pred = model(X_train)
        loss = criterion(y_pred, y_train)

        predicted = torch.max(y_pred.data, 1)[1]
        batch_correct = (predicted==y_train).sum()
        train_correct+=batch_correct

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        if batch%256 == 0:
            print(f"Epoch {i+1} Batch {batch} loss: {loss.item()}")

    train_losses.append(loss)
    train_corrects.append(train_correct)

    test_correct = 0
    with torch.no_grad():
        for batch, (X_test, y_test) in enumerate(test_loader):

```

```

X_test = X_test.to(device)
y_test = y_test.to(device)

y_val = model(X_test)

predicted = torch.max(y_val.data,1)[1]

test_correct+=(predicted==y_test).sum()

loss = criterion(y_val, y_test)
test_losses.append(loss)
test_corrects.append(test_correct)

current_time = time.time()
total = current_time - start_time
print(f"Training finished in {total/60} minutes")

for i in range(len(train_losses)):
    train_losses[i] = train_losses[i].cpu().detach().numpy()

for i in range(len(test_losses)):
    test_losses[i] = test_losses[i].cpu().detach().numpy()

print("train loss: ", train_losses)
print("test_loss : ", train_losses)

plt.plot(train_losses, label='train_loss')
plt.plot(test_losses, label='test_loss')
plt.ylabel('Losses')
plt.xlabel('epoch')
plt.xticks(range(0,5))
plt.legend()
print(plt.show())

torch.manual_seed(42)
test_loader_all = DataLoader(test_data, batch_size=10000,
shuffle=False)
test_loader_mini = DataLoader(test_data, batch_size=100,
shuffle=False)
preds = []

with torch.no_grad():
    correct = 0
    for batch, (X_test, y_test) in enumerate(test_loader_mini):
        batch+=1
        X_test = X_test.to(device)
        y_test = y_test.to(device)

        y_val = model(X_test)
        predicted = torch.max(y_val, 1)[1]
        preds.append(predicted)

        correct+=(predicted==y_test).sum()

y_test = test_data.targets
y_test = y_test.numpy()

y_pred = []

for batch_y in preds:
    for y_val in batch_y:

```

```

y_pred.append(y_val)

for i in range(len(y_pred)):
    y_pred[i] = y_pred[i].cpu().detach().numpy()

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))

```

- **Arsitektur Model Sebelum *Fine Tune***

```

AlexNet(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=4096, out_features=1000, bias=True)
  )
)

```

- **Arsitektur Model Setelah *Fine Tune* (Mengubah jumlah channel input, kelas output dan menambah layer logsoftmax)**

```

AlexNet(
  (features): Sequential(
    (0): Conv2d(1, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))
    (1): ReLU(inplace=True)
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (4): ReLU(inplace=True)
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (7): ReLU(inplace=True)
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (9): ReLU(inplace=True)
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (11): ReLU(inplace=True)
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))
  (classifier): Sequential(
    (0): Dropout(p=0.5, inplace=False)
    (1): Linear(in_features=9216, out_features=4096, bias=True)
    (2): ReLU(inplace=True)
    (3): Dropout(p=0.5, inplace=False)
    (4): Linear(in_features=4096, out_features=4096, bias=True)
    (5): ReLU(inplace=True)
    (6): Linear(in_features=4096, out_features=10, bias=True)
    (7): LogSoftmax(dim=1)
  )
)

```


- Freeze Layer (Hanya layer input dan output yang di update bobotnya / *require grad*)


```
parameter name features.0.weight gradient : True
parameter name features.0.bias gradient   : True
parameter name features.3.weight gradient : False
parameter name features.3.bias gradient   : False
parameter name features.6.weight gradient : False
parameter name features.6.bias gradient   : False
parameter name features.8.weight gradient : False
parameter name features.8.bias gradient   : False
parameter name features.10.weight gradient : False
parameter name features.10.bias gradient   : False
parameter name classifier.1.weight gradient : False
parameter name classifier.1.bias gradient   : False
parameter name classifier.4.weight gradient : False
parameter name classifier.4.bias gradient   : False
parameter name classifier.6.weight gradient : True
parameter name classifier.6.bias gradient   : True
```

- Proses Pelatihan

```
Epoch 1 Batch 256 loss: 0.5709512829780579
Epoch 2 Batch 256 loss: 0.271210640668869
Epoch 3 Batch 256 loss: 0.25392141938209534
Training finished in 6.8701050043106076 minutes
```

Waktu pelatihan: ± 7 Menit

- Hasil Loss Pelatihan



- Hasil Akurasi Terhadap Data *Test*

	precision	recall	f1-score	support
0	0.96	0.96	0.96	980
1	0.98	0.99	0.98	1135
2	0.91	0.94	0.92	1032
3	0.92	0.93	0.92	1010
4	0.95	0.93	0.94	982
5	0.90	0.95	0.93	892
6	0.92	0.95	0.93	958
7	0.93	0.95	0.94	1028
8	0.94	0.85	0.89	974
9	0.91	0.89	0.90	1009
accuracy			0.93	10000
macro avg	0.93	0.93	0.93	10000
weighted avg	0.93	0.93	0.93	10000

Hasil akurasi terhadap data uji : 93%

- Hasil *Confusion Matrix*

[[939	0	6	0	1	7	16	1	9	1]
[0	1118	4	2	0	0	8	1	2	0]	
[9	1	966	24	3	11	3	9	4	2]	
[0	0	28	940	0	25	0	5	8	4]	
[5	1	3	0	911	0	21	10	3	28]	
[5	2	7	15	1	845	3	2	8	4]	
[11	2	10	1	8	14	908	1	3	0]	
[0	6	17	6	8	3	1	975	0	12]	
[7	3	17	27	4	18	26	5	830	37]	
[6	4	2	10	24	12	2	38	16	895]]	